

# Semantic Role Labeling with Conditional Random Field and SpanBERT

Slides originally by Tuur Leeuwenberg and Quynh Do, adapted by Elias Moons, Sumam Francis, Ruiqi Li, Nathan Cornille, Ruben Cartuyvels, Marie-Francine Moens

**LIIR lab, Department of Computer Science, KU Leuven**

2025

# Semantic Role Labeling (SRL)

**Who** does **what** to **whom** and **how**?

He? loves<sub>*predicate*</sub> dogs? tremendously?

He? ran<sub>*predicate*</sub> fast?

John? dances<sub>*predicate*<sub>1</sub></sub> while Mary? sings<sub>*predicate*<sub>2</sub></sub>

# Semantic Role Labeling (SRL)

**Who** does **what** to **whom** and **how**?

He<sub>A0</sub> loves<sub>*predicate*</sub> dogs<sub>A1</sub> tremendously<sub>*AM-MNR*</sub>

He<sub>A0</sub> ran<sub>*predicate*</sub> fast<sub>*AM-MNR*</sub>

John<sub>A01</sub> dances<sub>*predicate1*</sub> while Mary<sub>A02</sub> sings<sub>*predicate2*</sub>

*A0* = agent, *A1* = patient, *AM – MNR* = manner.

# SRL as Sequence Prediction

Input sentence and predicate index  $(s, j)$ :

He? loves<sub>*predicate*</sub> dogs? tremendously?

Output label sequence:

$$L = [A0, A1, AM - MNR]$$

Model:

$$P(L|s, j)$$

## Sequence Prediction with a Linear Chain CRF

$P(L|s, j)$  is calculated using weights  $\lambda_k \in \lambda$  and features  $f_k \in F$  by first scoring the label configuration  $L$ .

$$score(L|s, j) = \sum_{i=1}^{C(s, j)} \sum_{k=1}^K \lambda_k f_k(s, i, j, l_i, l_{i-1})$$

with  $C(s, j)$  is the total number of arguments of the given predicate.

## Sequence Prediction with a Linear Chain CRF

$P(L|s, j)$  is calculated using weights  $\lambda_k \in \lambda$  and features  $f_k \in F$  by first scoring the label configuration  $L$ .

$$score(L|s, j) = \sum_{i=1}^{C(s,j)} \sum_{k=1}^K \lambda_k f_k(s, i, j, l_i, l_{i-1})$$

with  $C(s, j)$  is the total number of arguments of the given predicate.

The score is made a probability by normalizing over all possible labelings for the sequence (softmax):

$$P(L|s, j) = \frac{exp(score(L|s,j))}{\sum_{L'} exp(score(L'|s,j))}$$

## Feature Function

Each feature function  $f_k(s, i, j, l_i, l_{i-1})$  is a function taking as input:

- ① Sentence  $s$
- ② Position  $i$  of an argument in the sentence
- ③ Position  $j$  of the predicate
- ④ Label  $l_i$  of the current word
- ⑤ Label  $l_{i-1}$  of the previous argument.

and often outputs binary values (0 or 1).

## Example Feature: Argument word-form

For each possible (argument word, label) pair  $(w, y)$  in the training data, we have one feature function:

$$f_{w,y}(s, i, j, l_i, l_{i-1}) = \begin{cases} 1 & \text{if } a(s, j)[i] = w \text{ and } l_i = y \\ 0 & \text{otherwise} \end{cases}$$

## Example Feature: Argument word-form

For each possible (argument word, label) pair  $(w, y)$  in the training data, we have one feature function:

$$f_{w,y}(s, i, j, l_i, l_{i-1}) = \begin{cases} 1 & \text{if } a(s, j)[i] = w \text{ and } l_i = y \\ 0 & \text{otherwise} \end{cases}$$

Example Revisited:

He<sub>A0</sub> loves<sub>predicate</sub> dogs<sub>A1</sub> tremendously<sub>AM-MNR</sub>

He<sub>A0</sub> ran<sub>predicate</sub> fast<sub>AM-MNR</sub>

12 extracted features (4 word types, 3 label types):

(he,A0), (he,A1), (he,AM-MNR), (dogs,A0), (dogs,A1), (dogs,AM-MNR),  
(fast,A0), (fast,A1), (fast,AM-MNR), (tremendously,A0), (tremendously,  
A1), (tremendously,AM-MNR)

## Feature: Predicate argument-roles

For each possible (predicate word ( $p$ ), label ( $y$ )) pair in the training data, we have one feature function:

$$f_{p,y}(s, i, j, l_i, l_{i-1}) = \begin{cases} 1 & \text{if } s[j] = p \text{ and } l_i = y \\ 0 & \text{otherwise} \end{cases}$$

Example Revisited:

He<sub>A0</sub> loves<sub>*predicate*</sub> dogs<sub>A1</sub> tremendously<sub>*AM-MNR*</sub>

He<sub>A0</sub> ran<sub>*predicate*</sub> fast<sub>*AM-MNR*</sub>

6 extracted features (2 predicate types, 3 label types):

(ran,A0),(ran,A1),(ran,AM-MNR),(loves, A0), (loves,A1),  
(loves,AM-MNR)

## Feature: Label-transitions

For each possible (previous label ( $l_{i-1}$ ), current label ( $l_i$ )) pair in the training data, we have one feature function:

$$f_{l_1, l_2}(s, i, j, l_i, l_{i-1}) = \begin{cases} 1 & \text{if } l_i = l_2 \text{ and } l_{i-1} = l_1 \\ 0 & \text{otherwise} \end{cases}$$

Example Revisited:

He<sub>A0</sub> loves<sub>*predicate*</sub> dogs<sub>A1</sub> tremendously<sub>*AM-MNR*</sub>

He<sub>A0</sub> ran<sub>*predicate*</sub> fast<sub>*AM-MNR*</sub>

Extracted features: (*Null*, *A0*), (*Null*, *A1*), (*Null*, *AM – MNR*), (*A0*, *A1*), (*A1*, *A0*), (*A0*, *AM – MNR*), (*AM – MNR*, *A0*), (*A1*, *AM – MNR*), (*AM – MNR*, *A1*)

# Feature Engineering

- ① The combination of labels and instances' traditional features (e.g., argument word form, predicate word form) may result in a large number of parameters.
- ② But, many of these features never occur in the training data e.g., (he,AM-MNR).
- ③ Perhaps surprisingly, these features can be useful: they can be given a negative weight that prevents the spurious label from being assigned high probability.
- ④ Including these features typically results in slight improvements in accuracy, at the cost of greatly increasing the number of parameters in the model.

## Training: Objective function

Our objective is to find a set of parameters  $\lambda$  so that the resulting distribution  $P(L|s, j, \lambda)$  best fits the set of training examples.

A standard way of training CRFs is by **maximum log-likelihood**:

$$l(\lambda) = \log\left(\prod_{s,j} P(L|s, j)\right) = \sum_{s,j} \log(P(L|s, j)) =$$
$$\sum_{s,j} \log\left(\frac{\exp\left(\sum_{i=1}^{C(s,j)} \sum_{k=1}^K \lambda_k f_k(s, i, j, l_i, l_{i-1})\right)}{\sum_{L'} \exp\left(\sum_{i=1}^{C(s,j)} \sum_{k=1}^K \lambda_k f_k(s, i, j, y'_i, y'_{i-1})\right)}\right)$$

## Training: Stochastic Gradient Descent

For each sentence  $s$ , predicate position  $j$ , and the gold labeling  $L$ , we compute the gradient of the local likelihood  $l^{s,j}(\lambda) = \log(P(L|s, j))$  with respect to each feature weight  $\lambda_q$ :

$$\frac{\partial}{\partial \lambda_q} \log P(L|s, j) =$$
$$\sum_{i=1}^{C(s,j)} f_q(s, i, j, l_i, l_{i-1}) - \sum_{L'} [P(L'|s, j) * \sum_{i=1}^{C(s,j)} f_q(s, i, j, y'_i, y'_{i-1})]$$

# Deriving the gradient

$$\begin{aligned} & \frac{\partial}{\partial \lambda_q} \log P(L|s, j) \\ &= \frac{\partial}{\partial \lambda_q} \log \left( \frac{\exp(\sum_{i=1}^{C(s,j)} \sum_{k=1}^K \lambda_k f_k(s, i, j, l_i, l_{i-1}))}{\sum_{L'} \exp(\sum_{i=1}^{C(s,j)} \sum_{k=1}^K \lambda_k f_k(s, i, j, l'_i, l'_{i-1}))} \right) \\ &= \frac{\partial}{\partial \lambda_q} \log \left( \exp \left( \sum_{i=1}^{C(s,j)} \sum_{k=1}^K \lambda_k f_k(s, i, j, l_i, l_{i-1}) \right) \right) - \frac{\partial}{\partial \lambda_q} \log \left( \sum_{L'} \exp \left( \sum_{i=1}^{C(s,j)} \sum_{k=1}^K \lambda_k f_k(s, i, j, l'_i, l'_{i-1}) \right) \right) \\ &= \frac{1}{\exp(\sum_{i=1}^{C(s,j)} \sum_{k=1}^K \lambda_k f_k(s, i, j, l_i, l_{i-1}))} \exp \left( \sum_{i=1}^{C(s,j)} \sum_{k=1}^K \lambda_k f_k(s, i, j, l_i, l_{i-1}) \right) * \sum_{i=1}^{C(s,j)} f_q(s, i, j, l_i, l_{i-1}) \\ &\quad - \frac{1}{\sum_{L''} \exp(\sum_{i=1}^{C(s,j)} \sum_{k=1}^K \lambda_k f_k(s, i, j, l''_i, l''_{i-1}))} \sum_{L'} \left[ \exp \left( \sum_{i=1}^{C(s,j)} \sum_{k=1}^K \lambda_k f_k(s, i, j, l'_i, l'_{i-1}) \right) * \sum_{i=1}^{C(s,j)} f_q(s, i, j, l'_i, l'_{i-1}) \right] \\ &= \sum_{i=1}^{C(s,j)} f_q(s, i, j, l_i, l_{i-1}) - \sum_{L'} \left[ \frac{\exp(\sum_{i=1}^{C(s,j)} \sum_{k=1}^K \lambda_k f_k(s, i, j, l'_i, l'_{i-1}))}{\sum_{L''} \exp(\sum_{i=1}^{C(s,j)} \sum_{k=1}^K \lambda_k f_k(s, i, j, l''_i, l''_{i-1}))} * \sum_{i=1}^{C(s,j)} f_q(s, i, j, l'_i, l'_{i-1}) \right] \\ &= \sum_{i=1}^{C(s,j)} f_q(s, i, j, l_i, l_{i-1}) - \sum_{L'} P(L'|s, j) * \sum_{i=1}^{C(s,j)} f_q(s, i, j, l'_i, l'_{i-1}) \end{aligned}$$

# Stochastic Gradient Descent

Update the weights using the computed gradient:

$$\lambda_q = \lambda_q + \alpha * \frac{\partial}{\partial \lambda_q} \log P(L|s, j)$$

with  $\alpha$  a learning rate.

## SRL - Training<sup>1</sup> [Exercise\*]

We consider only the following sentence as all training data:

$\text{He}_{A0} \text{ ran}_{predicate} \text{ fast}_{AM-MNR}$

Using this training data, the feature set has 6 observed features:

- Argument word-forms:  $(he, A0), (fast, AM - MNR)$
- Predicate argument-roles:  $(ran, A0), (ran, AM - MNR)$
- Label transitions:  $(Null, A0), (A0, AM - MNR)$

So a feature vector is constructed by evaluating the following features for each (argument, predicate) pair:

$[f_{he,A0}, f_{fast,AM-MNR}, f_{ran,A0}, f_{ran,AM-MNR}, f_{Null,A0}, f_{A0,AM-MNR}]$

We have two (argument, predicate) pairs:

- ① (ran,he) with label A0 results in  $[1, 0, 1, 0, 1, 0]$
- ② (ran,fast) with label AM-MNR gives  $[0, 1, 0, 1, 0, 1]$

<sup>1</sup>to make it simpler for the exercise, we only consider features that are observed in the training data (at least one training instance has value of 1).

## SRL - Training

At the start of training we set our feature weights  $\lambda$  to zero.

Initialized weights:  $\lambda = [0.0, 0.0, 0.0, 0.0, 0.0, 0.0]$

## Iteration 1

$\text{He}_{A0} \text{ ran}_{predicate} \text{ fast}_{AM-MNR}$

From the above training example we obtain gold labeling  $L$

$$L = [A0, AM - MNR]$$

For the example we have (only 2) possible label assignments  $L'$

$$L' \in \{[A0, AM - MNR], [AM - MNR, A0]\}.$$

To update  $\lambda$  with SGD we set learning rate  $\alpha = 0.1$  and calculate gradients using the gold labeling.

$$\lambda_{new} = \lambda_{old} + 0.1 * \frac{\partial}{\partial \lambda} logP([A0, AM - MNR] | s, 2)$$

## Calculating gradients $\log P([A0, AM\text{-}MNR]|s, 2)$

$$\frac{\partial}{\partial \lambda_q} \log P([A0, AM\text{-}MNR]|s, j) =$$

$$\sum_{i=1}^{C(s,j)} f_q(s, i, j, l_i, l_{i-1}) - \sum_{L'} [P(L'|s, j) * \sum_{i=1}^{C(s,j)} f_q(s, i, j, y'_i, y'_{i-1})]$$

$$= f(s, 1, 2, A0, Null) + f(s, 3, 2, AM\text{-}MNR, A0) -$$

$$P([A0, AM\text{-}MNR]|s, 2) * (f(s, 1, 2, A0, Null) + f(s, 3, 2, AM\text{-}MNR, A0)) -$$

$$P([AM\text{-}MNR, A0]|s, 2) * (f(s, 1, 2, AM\text{-}MNR, Null) + f(s, 3, 2, A0, AM\text{-}MNR))$$

$$= [1, 0, 1, 0, 1, 0] + [0, 1, 0, 1, 0, 1] -$$

$$1/2 * ([1, 0, 1, 0, 1, 0] + [0, 1, 0, 1, 0, 1]) -$$

$$1/2 * ([0, 0, 0, 1, 0, 0] + [0, 0, 1, 0, 0, 0])$$

$$= [0.5, 0.5, 0, 0, 0.5, 0.5]$$

where

$$\begin{aligned} P([A0, AM\text{-}MNR]|s, 2) &= \\ \frac{\exp(score([A0, AM\text{-}MNR]|s, 2))}{\exp(score([A0, AM\text{-}MNR]|s, 2)) + \exp(score([AM\text{-}MNR, A0]|s, 2))} \\ &= \frac{\exp(0 + 0)}{\exp(0 + 0) + \exp(0 + 0)} = 1/2 \end{aligned}$$

where

$$\begin{aligned} score([A0, AM\text{-}MNR]|s, 2) &= \sum_{i=1}^{C(s,j)} \sum_{k=1}^K \lambda_k f_k(s, i, j, l_i, l_{i-1}) \\ &= \lambda * f(s, 1, 2, A0, Null) + \lambda * f(s, 3, 2, AM\text{-}AMNR, A0) = 0 \\ &= ([0, 0, 0, 0, 0, 0] * [1, 0, 1, 0, 1, 0]) + ([0, 0, 0, 0, 0, 0] * [0, 1, 0, 1, 0, 1]) \\ &= 0 \end{aligned}$$

and

$$score([AM\text{-}MNR, A0]|s, 2) = \lambda * f(s, 1, 2, AM\text{-}MNR, Null) + \lambda * f(s, 3, 2, A0, AM\text{-}MNR) = 0$$

Similarly, we also obtain  $P([AM\text{-}MNR, A0]|s, 2) = 1/2$

Update weights:

$$\lambda_{new} = \lambda_{old} + 0.1 * [0.5, 0.5, 0, 0, 0.5, 0.5] = [0.05, 0.05, 0, 0, 0.05, 0.05]$$

## Iteration 2 (Same procedure but with a different $\lambda$ )

Exactly the same except that  $P([A0, AM-MNR]|s, 2)$  and  $P([A0, AM-MNR]|s, 2)$  are different because  $\lambda$  is different.

$$\begin{aligned}\frac{\partial}{\partial \lambda} \log P((A0, AM - MNR)|s, 2) &= \\ &= f(s, 1, 2, A0, Null) + f(s, 3, 2, AM-MNR, A0) - \\ &\quad P([A0, AM-MNR]|s, 2) * (f(s, 1, 2, A0, Null) + f(s, 3, 2, AM - MNR, A0)) - \\ &\quad P([AM-MNR, A0]|s, 2) * (f(s, 1, 2, AM-MNR, Null) + f(s, 3, 2, A0, AM-MNR)) \\ &= [1, 0, 1, 0, 1, 0] + [0, 1, 0, 1, 0, 1] - \\ &\quad .55 * ([1, 0, 1, 0, 1, 0] + [0, 1, 0, 1, 0, 1]) - \\ &\quad .45 * ([0, 0, 0, 1, 0, 0] + [0, 0, 1, 0, 0, 0]) \\ &= [0.45, 0.45, 0, 0, 0.45, 0.45]\end{aligned}$$

Update weights:

$$\begin{aligned}\lambda_{new} &= \lambda_{old} + 0.1 * [0.45, 0.45, 0, 0, 0.45, 0.45] = [0.05, 0.05, 0, 0, 0.05, 0.05] + \\ &[0.045, 0.045, 0, 0, 0.045, 0.045] = [0.095, 0.095, 0, 0, 0.095, 0.095]\end{aligned}$$

where

$$\begin{aligned} P([A0, AM-MNR]|s, 2) &= \\ \frac{\exp(score([A0, AM-MNR]|s, 2))}{\exp(score([A0, AM-MNR]|s, 2)) + \exp(score([AM-MNR, A0]|s, 2))} \\ &= \frac{\exp(0.2)}{\exp(0.2) + \exp(0)} \approx .55 \end{aligned}$$

where

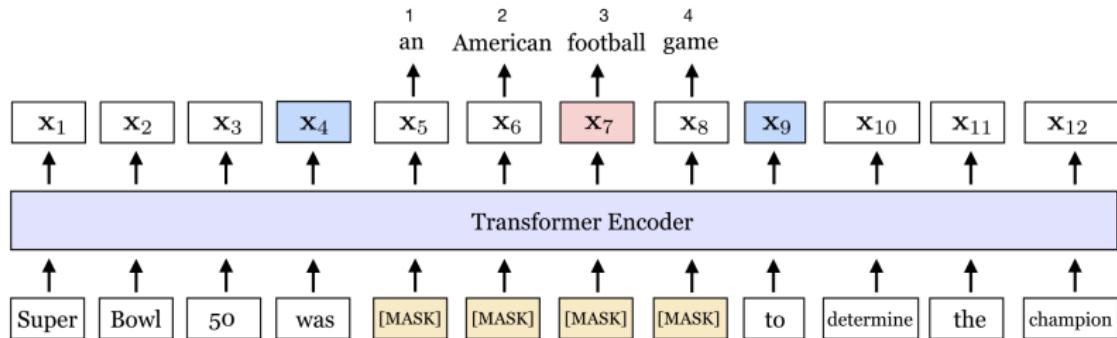
$$\begin{aligned} score([A0, AM-MNR]|s, 2) &= \sum_{i=1}^{C(s,j)} \sum_{k=1}^K \lambda_k f_k(s, i, j, l_i, l_{i-1}) \\ &= \lambda * f(s, 1, 2, A0, Null) + \lambda * f(s, 3, 2, AM-AMNR, A0) \\ &= ([0.05, 0.05, 0, 0, 0.05, 0.05] * [1, 0, 1, 0, 1, 0]) \\ &\quad + ([0.05, 0.05, 0, 0, 0.05, 0.05] * [0, 1, 0, 1, 0, 1]) \\ &\approx .2 \end{aligned}$$

and

$$\begin{aligned} score([AM-MNR, A0]|s, 2) &= \lambda * f(s, 1, 2, AM-MNR, Null) + \lambda * f(s, 3, 2, A0, AM-MNR) \\ &= ([0.05, 0.05, 0, 0, 0.05, 0.05] * [0, 0, 0, 1, 0, 0]) \\ &\quad + ([0.05, 0.05, 0, 0, 0.05, 0.05] * [0, 0, 1, 0, 0, 0]) \\ &= 0 \end{aligned}$$

# SpanBERT: pre-training

$$\begin{aligned}\mathcal{L}(\text{football}) &= \mathcal{L}_{\text{MLM}}(\text{football}) + \mathcal{L}_{\text{SBO}}(\text{football}) \\ &= -\log P(\text{football} \mid \mathbf{x}_7) - \log P(\text{football} \mid \mathbf{x}_4, \mathbf{x}_9, \mathbf{p}_3)\end{aligned}$$



The span *an American football game* is masked.

$x_4$  and  $x_9$  are the output representations of the boundary tokens.

SBO predicts each token in the masked span.

$p_3$  is the position embedding of the third token of the span.

## SpanBERT pre-training: Spans to mask

- Random sub-word tokens
- Random whole words
- Noun Phrases
- Named Entities

## SpanBERT: downstream SRL

Does this make sense?

- Arguments usually cover a span of words, so yes.
- SpanBERT might be overly complex for SRL?
- Is there an advantage of using these span representations instead of predicting per individual word what SRL argument it belongs to?
- How do you determine span boundaries, or do you test all possible spans (there are  $\mathcal{O}(T^2)$  possible spans in a length  $T$  sentence)?

# SpanBERT: Objective for downstream SRL

Sentences  $X^n, n \in 1 \dots N$  as sequences of  $T^n$  words  $x_1^n \dots x_{T^{(n)}}^n$ .

Each sentence has a set of  $S^n$  masked spans. Each span is defined by its length  $K_i^n$  and start position  $c_i^n$ .

Each word  $x_j^n$  has a position embedding  $p_j^n$  and an SRL label  $z_j^n$ .

$\Theta_c$  are the common model parameters involved in both objectives,  $\Theta_{SBO}$  the parameters only involved in the SBO loss, and  $\Theta_{CAT}$  those only involved in the CAT loss.

Then:

$$\mathcal{L} = \sum_{n=1}^N \mathcal{L}^n$$

$$\mathcal{L}^n =$$

# SpanBERT: Objective for SRL

Sentences  $X^n, n \in 1 \dots N$  as sequences of  $T^n$  words  $x_1^n \dots x_{T(n)}^n$ .

Each sentence has a set of  $S^n$  masked spans. Each span is defined by its length  $K_i^n$  and start position  $c_i^n$ .

Each word  $x_j^n$  has a position embedding  $p_j^n$  and an SRL label  $z_j^n$ .

$\Theta_c$  are the common model parameters involved in both objectives,  $\Theta_{SBO}$  the parameters only involved in the SBO loss, and  $\Theta_{CAT}$  those only involved in the CAT loss.

Then:

$$\begin{aligned}\mathcal{L} &= \sum_{n=1}^N \mathcal{L}^n \\ \mathcal{L}^n &= \sum_{i=1}^{S^n} \sum_{k=0}^{K_i^n - 1} \mathcal{L}_{SBO}(x_{c_i^n+k}^n) + \sum_{t=1}^{T^n} \mathcal{L}_{CAT}(z_t^n) \\ &= \sum_{i=1}^{S^n} \sum_{k=0}^{K_i^n - 1} \left[ -\log P(x_{c_i^n+k}^n | x_{c_i^n-1}^n, x_{c_i^n+K_i^n}^n, p_{c_i^n+k}^n; \Theta_c, \Theta_{SBO}) \right] \\ &\quad + \sum_{t=1}^{T^n} -\log P(z_t^n | X^n; \Theta_c, \Theta_{CAT})\end{aligned}$$